


# Chapter 4

## Basic SQL



Sixth Edition

# Fundamentals of Database Systems

Elmasri • Navathe

Addison-Wesley  
is an imprint of

PEARSON

# Chapter 4

## Basic SQL

# Chapter 4 Outline

- Overview of SQL
- SQL Data Definition (DDL) for Specifying a Relational Database Schema
  - CREATE statements
  - Specifying Constraints
- Basic Retrieval Queries in SQL
  - SELECT ... FROM ... WHERE ... statements
- Basic Database Modification in SQL
  - INSERT, DELETE, UPDATE statements

# Overview of SQL

- SQL is a standard, comprehensive language, based on the relational model
- SQL includes capabilities for many functions:
  - DDL statements for creating schemas and specifying data types and constraints
  - Statements for specifying database retrievals
  - Statements for modifying the database
  - Statements for specifying views, triggers, and assertions (see Chapter 5)
  - Many other capabilities
- We introduce the first three capabilities in this chapter

# SQL Features in This Chapter

- Basic SQL DDL
  - Includes the CREATE statements
  - Has a comprehensive set of SQL data types
  - Can specify key, referential integrity, and other constraints
  - Object-oriented features of SQL DDL will be presented in Chapter 11
- Basic Retrieval Queries in SQL
  - SELECT ... FROM ... WHERE ... statements
- Basic Database Modification in SQL
  - INSERT, DELETE, UPDATE statements

# SQL: Origins and History

- Evolution of the SQL Standard
  - First standard approved in 1989 (called SQL-89 or SQL1)
  - Comprehensive revision in 1992 (SQL-92 or SQL2)
  - Third big revision in 1999 (SQL-99, a trimmed-down version of a more ambitious revision known as SQL3)
  - Other revisions known as SQL:2003, SQL:2006, SQL:2008
- Origins of SQL
  - Originally called SEQUEL (Structured English Query Language), then SQL (Structured Query Language)
  - Developed at IBM Research for experimental relational DBMS called *System-R* in the 1970s

# SQL Data Definition

- CREATE statement can be used to:
  - Create a named database schema
  - Create individual database tables and specify data types for the table attributes, as well as *key*, *referential integrity*, and *NOT NULL* constraints
  - Create named constraints
- Other commands can modify the tables and constraints
  - DROP and ALTER statements (discussed in Chapter 5)

# CREATE TABLE

- In its *simplest form*, specifies a new base relation by giving it a name, and specifying each of its attributes and their data types (INTEGER, FLOAT, DECIMAL(i,j), CHAR(n), VARCHAR(n), DATE, and other data types)
- A constraint NOT NULL may be specified on an attribute

```
CREATE TABLE DEPARTMENT (  
  DNAME      VARCHAR(15)  NOT NULL,  
  DNUMBER    INT          NOT NULL,  
  MGRSSN     CHAR(9)      NOT NULL,  
  MGRSTARTDATE DATE      ) ;
```



# CREATE TABLE (cont.)

- CREATE TABLE can also specify the primary key, UNIQUE keys, and referential integrity constraints (foreign keys)
- Via the PRIMARY KEY, UNIQUE, and FOREIGN KEY phrases

```
CREATE TABLE DEPARTMENT (  
    DNAME          VARCHAR(15)  NOT NULL,  
    DNUMBER        INT           NOT NULL,  
    MGRSSN         CHAR(9)      NOT NULL,  
    MGRSTARTDATE   DATE,  
    PRIMARY KEY (DNUMBER) ,  
    UNIQUE (DNAME) ,  
    FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE) ;
```

# Example: The COMPANY Database

- Figure 3.7 (next slide) shows the COMPANY database *schema diagram* introduced in Chapter 3
- Referential integrity constraints shown as directed edges from *foreign key* to *referenced relation*
- Primary key attributes of each table *underlined*

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

**Figure 3.7**  
Referential integrity constraints displayed on the COMPANY relational database schema.

# Example: The COMPANY DDL

- Figure 4.1 (following two slides) shows *example DDL* for creating the tables of the COMPANY database
- Circular reference problem:
  - In Figure 3.7, some foreign keys cause *circular references*:
    - EMPLOYEE.Dno → DEPARTMENT.Dnumber
    - DEPARTMENT.Mgr\_ssn → EMPLOYEE.Ssn
  - One of the tables must be created first without the FOREIGN KEY in the CREATE TABLE statement
    - The missing FOREIGN KEY can be added later using ALTER TABLE (see Chapter 5)

## CREATE TABLE EMPLOYEE

( Fname	VARCHAR(15)	NOT NULL,
Minit	CHAR,	
Lname	VARCHAR(15)	NOT NULL,
Ssn	CHAR(9)	NOT NULL,
Bdate	DATE,	
Address	VARCHAR(30),	
Sex	CHAR,	
Salary	DECIMAL(10,2),	
Super_ssn	CHAR(9),	
Dno	INT	NOT NULL,

**PRIMARY KEY** (Ssn),

**FOREIGN KEY** (Super\_ssn) **REFERENCES** EMPLOYEE(Ssn),

**FOREIGN KEY** (Dno) **REFERENCES** DEPARTMENT(Dnumber) );

## CREATE TABLE DEPARTMENT

( Dname	VARCHAR(15)	NOT NULL,
Dnumber	INT	NOT NULL,
Mgr_ssn	CHAR(9)	NOT NULL,
Mgr_start_date	DATE,	

**PRIMARY KEY** (Dnumber),

**UNIQUE** (Dname),

**FOREIGN KEY** (Mgr\_ssn) **REFERENCES** EMPLOYEE(Ssn) );

## CREATE TABLE DEPT\_LOCATIONS

( Dnumber	INT	NOT NULL,
Dlocation	VARCHAR(15)	NOT NULL,

**PRIMARY KEY** (Dnumber, Dlocation),

**FOREIGN KEY** (Dnumber) **REFERENCES** DEPARTMENT(Dnumber) );

### Figure 4.1

SQL CREATE TABLE data definition statements for defining the COMPANY schema from Figure 3.7.

## CREATE TABLE PROJECT

( Pname	VARCHAR(15)	NOT NULL,
Pnumber	INT	NOT NULL,
Plocation	VARCHAR(15),	
Dnum	INT	NOT NULL,

**PRIMARY KEY** (Pnumber),

**UNIQUE** (Pname),

**FOREIGN KEY** (Dnum) **REFERENCES** DEPARTMENT(Dnumber) );

## CREATE TABLE WORKS\_ON

( Essn	CHAR(9)	NOT NULL,
Pno	INT	NOT NULL,
Hours	DECIMAL(3,1)	NOT NULL,

**PRIMARY KEY** (Essn, Pno),

**FOREIGN KEY** (Essn) **REFERENCES** EMPLOYEE(Ssn),

**FOREIGN KEY** (Pno) **REFERENCES** PROJECT(Pnumber) );

## CREATE TABLE DEPENDENT

( Essn	CHAR(9)	NOT NULL,
Dependent_name	VARCHAR(15)	NOT NULL,
Sex	CHAR,	
Bdate	DATE,	
Relationship	VARCHAR(8),	

**PRIMARY KEY** (Essn, Dependent\_name),

**FOREIGN KEY** (Essn) **REFERENCES** EMPLOYEE(Ssn) );

# The CREATE SCHEMA Statement

- A DBMS can manage multiple databases
- DBA (or authorized users) can use CREATE SCHEMA to have distinct databases; for example:

```
CREATE SCHEMA COMPANY AUTHORIZATION 'Smith';
```

- Each database has a schema name (e.g. COMPANY)
- User 'Smith' will be “owner” of schema, can create tables, other constructs in that schema
- Table names can be prefixed by schema name if multiple schemas exist (e.g. COMPANY.EMPLOYEE)
- Prefix not needed if there is a “default schema” for the user

# Basic SQL Data Types

- **Basic numeric data types:**

- Integers: INTEGER (or INT), SMALLINT
- Real (floating point): FLOAT (or REAL), DOUBLE PRECISION
- Formatted: DECIMAL(i,j) (or DEC (i,j) or NUMERIC(i,j)) specify i total decimal digits, j after decimal point
  - i called *precision*, j called *scale*

- **Basic character string data types:**

- Fixed-length: CHAR(n) or CHARACTER(n)
- Variable-length: VARCHAR(n) or CHAR VARYING(n) or CHARACTER VARYING(n)
- Placed in single quotation marks, case sensitive
- Fixed length strings padded with blanks to the right
- Alphabetically ordered
- Concatenation operator ||

CLOB



# SQL Data Types (cont)

- Bit – string data types:
  - BIT(n), BIT VARYING (n), default length of n is 1
  - Literal bit strings – B'0101'
  - Variable length bit strings – BLOB in Kb/Mb/Gb
- Boolean data types:
  - TRUE/FALSE/UNKNOWN
- Large object data types:
  - Binary large objects: BLOB(n)
  - Can be used to store attributes that represent images, audio, video, or other large binary objects
  - Character large objects: CLOB(n)
  - Can be used attributes that store articles, news stories, text files, and other large character objects

# SQL Data Types (cont)

- DATE data type:
  - Standard DATE formatted as yyyy-mm-dd
  - For example, DATE '2010-01-22'
  - Older formats still used in some systems, such as 'JAN-22-2010'
  - Values are ordered, with later dates larger than earlier ones
- TIME data type:
  - Standard TIME formatted as hh:mm:ss
  - E.g., TIME '15:20:22' represents 3.20pm and 22 seconds
  - TIME(i) includes and additional i-1 decimal digits for fractions of a second
  - E.g., TIME(5) value could be '15:20:22.1234'
  - < comparison works for DATE and TIME
  - Literal values – DATE'2008-09-27' or TIME'09:10:12'

# SQL Data Types (cont)

- **TIMESTAMP** data type:
  - A DATE combined with a TIME(i), where the default i=7
  - For example, `TIMESTAMP '2010-01-22 15:20:22.123456'`
  - A different i>7 can be specified if needed
- **INTERVAL**
  - represents a relative YEAR/MONTH or DAY/TIME (e.g. 2 years and 5 months, or 5 days 20 hours 3 minutes 22 seconds)
  - Can be used to increment/decrement an absolute value of DATE/TIME/TIMESTAMP
- Other SQL data types exist; we presented the most common ones

# Specifying Constraints in SQL

Basic constraints in SQL as part of table creation (Key constraints, Referential integrity constraints, Restrictions on attribute domains & NULLs)

Attribute constraints & attribute defaults

- NOT NULL constraints, implicit for key attributes, can be defined for others.
- Default value for an attribute –
- Restrict attribute or domain value using CHECK

Key and Referential Integrity constraints

- PRIMARY key – PRIMARY KEY clause
- UNIQUE key –
- FOREIGN KEY clause

# Specifying Referential Integrity Options and Naming of Constraints

- We can specify RESTRICT (the default), CASCADE, SET NULL or SET DEFAULT on referential integrity constraints (foreign keys)
- Separate options for ON DELETE, ON UPDATE
- Figure 4.2 (next slide) gives some examples
- A constraint can be given a *constraint name*; this allows to DROP the constraint later (see Chapter 5)
- Constraint names should be unique in a schema
- Figure 4.2 illustrates naming of constraints

```

CREATE TABLE EMPLOYEE
(
    ...,
    Dno          INT          NOT NULL          DEFAULT 1,
    CONSTRAINT EMPPK
        PRIMARY KEY (Ssn),
    CONSTRAINT EMPSUPERFK
        FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
            ON DELETE SET NULL          ON UPDATE CASCADE,
    CONSTRAINT EMPDEPTFK
        FOREIGN KEY(Dno) REFERENCES DEPARTMENT(Dnumber)
            ON DELETE SET DEFAULT       ON UPDATE CASCADE);

CREATE TABLE DEPARTMENT
(
    ...,
    Mgr_ssn      CHAR(9)      NOT NULL          DEFAULT '888665555',
    ...,
    CONSTRAINT DEPTPK
        PRIMARY KEY(Dnumber),
    CONSTRAINT DEPTSK
        UNIQUE (Dname),
    CONSTRAINT DEPTMGRFK
        FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
            ON DELETE SET DEFAULT       ON UPDATE CASCADE);

CREATE TABLE DEPT_LOCATIONS
(
    ...,
    PRIMARY KEY (Dnumber, Dlocation),
    FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
        ON DELETE CASCADE              ON UPDATE CASCADE);

```

### Figure 4.2

Example illustrating how default attribute values and referential integrity triggered actions are specified in SQL.

# Specifying DOMAIN Constraints Using the CHECK Clause

- Another type of constraint can be specified using CHECK
- E.g., CHECK (Dnumber > 0 AND Dnumber < 21) can be specified after the Dnumber attribute
- Alternatively, a special domain can be created with a domain name using CREATE DOMAIN
- E.g. CREATE DOMAIN D\_NUM AS INTEGER CHECK (Dnumber > 0 AND Dnumber < 21);
- D\_NUM can now be used as the data type for the Dnumber attribute of the DEPARTMENT table, as well as for Dnum of PROJECT, Dno of EMPLOYEE, and so on
- CHECK can also be used to specify more general constraints (see Chapter 5)



# Basic Retrieval Queries in SQL

- SQL basic statement for retrieving information from a database is the **SELECT** statement
  - NOTE: This is *not the same* as the SELECT operation of the relational algebra (see Chapter 6)
- Important distinction between *practical SQL model* and the *formal relational model* (see Chapter 3):
  - SQL allows a table (relation) to have two or more tuples that are identical in all their attribute values
  - Hence, an SQL relation (table) is a **multi-set** (sometimes called a **bag**) of tuples; it is *not* a set of tuples
- SQL relations can be *constrained to be sets* by specifying PRIMARY KEY or UNIQUE attributes, or by using the DISTINCT option in a query



# Bag (Multiset) versus Set

- A **bag** or **multi-set** is like a set, but an element may appear more than once
  - Example:  $\{A, B, C, A\}$  is a bag (but *not* a set).  $\{A, B, C\}$  is a bag (and also a set).
  - Bags also resemble lists, but the order is irrelevant in a bag.
- Example:
  - $\{A, B, A\} = \{B, A, A\}$  as bags
  - However,  $[A, B, A]$  is not equal to  $[B, A, A]$  as lists

# Basic Retrieval Queries in SQL (cont.)

- Simplest form of the SQL SELECT statement is called a *mapping* or a SELECT-FROM-WHERE *block*
- Our examples use the COMPANY database schema
- **SELECT** <attribute list>  
**FROM** <table list>  
**WHERE** <condition>
  - <attribute list> is a list of attribute names whose values are to be retrieved by the query, also called *projection attributes*
  - <table list> is a list of the table (relation) names required to process the query
  - <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query, also known as *selection condition*

# Simple SQL Queries (cont.)

- Query text ends with a semi-colon
- First example of a simple query on one table (relation)
- Query 0: Retrieve the birthdate and address of the employee whose name is 'John B. Smith'
- Use the EMPLOYEE table only

```
Q0:SELECT  BDATE, ADDRESS  
      FROM    EMPLOYEE  
      WHERE   FNAME='John' AND MINIT='B'  
             AND LNAME='Smith' ;
```

# Simple SQL Queries (cont.)

- Example of a query that uses two tables
- Query 1: Retrieve the name and address of all employees who work for the 'Research' department.

```
Q1:SELECT  FNAME, LNAME, ADDRESS  
      FROM    EMPLOYEE, DEPARTMENT  
      WHERE   DNAME='Research' AND DNUMBER=DNO ;
```

- (DNAME='Research') is called a *selection condition*
- (DNUMBER=DNO) is called a *join condition* (it *joins* two tuples from EMPLOYEE and DEPARTMENT tables whenever EMPLOYEE.DNO=DEPARTMENT.DNUMBER)
- *Select-project-join* query

# Simple SQL Queries (cont.)

- A *selection condition* refers to attributes from a single table, and selects (chooses) those records that satisfy the condition
- A *join condition generally* refers to attributes from two tables, and joins (or combines) pairs of records that satisfy the condition
- In the previous query:
  - (DNAME='Research') chooses the DEPARTMENT record
  - (DNUMBER=DNO) joins the record with each EMPLOYEE record that works for that department

# Simple SQL Queries (cont.)

- Query 2 (needs 3 tables): For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
Q2: SELECT  PNUMBER, DNUM, LNAME, BDATE, ADDRESS  
FROM      PROJECT, DEPARTMENT, EMPLOYEE  
WHERE     DNUM=DNUMBER AND MGRSSN=SSN  
AND PLOCATION='Stafford' ;
```

- In Q2, there are two join conditions
- The join condition DNUM=DNUMBER relates a PROJECT record to its controlling DEPARTMENT record
- The join condition MGRSSN=SSN relates the controlling DEPARTMENT to the EMPLOYEE who manages that department

# Qualifying Attribute Name with Table Name

- An attribute name in an SQL query can be *prefixed* (or *qualified*) with the table name
- Examples:

**EMPLOYEE.LNAME**

**DEPARTMENT.DNAME**

Query Q1 can be rewritten as:

```
SELECT  EMPLOYEE.FNAME, EMPLOYEE.LNAME,  
          EMPLOYEE. ADDRESS
```

```
FROM      EMPLOYEE, DEPARTMENT
```

```
WHERE  DEPARTMENT.DNAME='Research' AND  
        DEPARTMENT.DNUMBER=EMPLOYEE.DNO ;
```

# Aliasing Table Names Using Tuple Variables

- An *alias* (or *tuple variable*) can be used instead of the table name when prefixing attribute names
- Example:

Query Q1 can be rewritten as follows using the aliases D for DEPARTMENT and E for EMPLOYEE:

```
SELECT  E.FNAME, E.LNAME, E.ADDRESS
FROM      EMPLOYEE AS E, DEPARTMENT AS D
WHERE     D.DNAME='Research' AND D.DNUMBER=E.DNO ;
```



# Renaming of Attributes

- It is also possible to **rename** the attributes of a table within a query; new attribute names are listed in the same order that the attributes were specified in CREATE TABLE
- This is sometimes necessary for certain joins (see Chapter 5)
- Example: Query Q1 can be rewritten as follows:

```
SELECT  E.FN, E.LN, E.ADR
FROM    DEPARTMENT AS D(DNM, DNO, MSSN, STRDATE),
          EMPLOYEE AS E(FN,MI,LN,S,BD,ADR,SX,SAL,SU,DNO)
WHERE    D.DNM='Research' AND D.DNO=E.DNO ;
```

# Queries that Require Qualifying of Attributes Names

- In SQL, all attribute names in a particular table must be different
  - However, the same attribute name can be used in *different tables*
  - In this case, it is required to use aliases if both tables are used in the same query
  - Example: **Suppose** that the attribute name NAME was used for both DEPARTMENT name and PROJECT name

Query: For each project, retrieve the project's name, and the name of its controlling department.

```
Q:  SELECT  P.NAME, D.NAME
     FROM    PROJECT AS P, DEPARTMENT AS D
     WHERE   P.DNUM=D.DNUMBER ;
```

In Q, P.NAME refers to the NAME attribute in the PROJECT table, and D.NAME refers to the NAME attribute in the DEPARTMENT table,

# Queries that Require Qualifying of Attributes (cont.)

- Some queries need to refer to the same relation twice
  - In this case, *aliases* are also required

Query 8: For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

```
Q8: SELECT  E.FNAME, E.LNAME, S.FNAME, S.LNAME
      FROM    EMPLOYEE AS E , EMPLOYEE AS S
      WHERE   E.SUPERSSN=S.SSN;
```

- In Q8, E and S are two different *aliases* or *tuple variables* for the EMPLOYEE relation
- We can think of E and S as two different *copies* of EMPLOYEE; E represents employees in role of *supervisees* and S represents employees in role of *supervisors*
- The join condition joins two different employee records together (a supervisor S and a subordinate E)

# Missing WHERE-clause

- The *WHERE-clause* is **optional** in an SQL query
- A *missing WHERE-clause* indicates no condition; hence, all tuples of the relations in the FROM-clause are selected
  - This is equivalent to the condition WHERE TRUE
- Example: Retrieve the SSN values for all employees.

```
Q9:  SELECT  SSN  
      FROM    EMPLOYEE ;
```

- If more than one relation is specified in the FROM-clause *and* there is no WHERE-clause (hence no join conditions), then *all possible combinations* of tuples are joined together (known as *CARTESIAN PRODUCT* of the relations)

# Missing WHERE-clause (cont.)

- Example:

```
Q10: SELECT SSN, DNAME  
      FROM EMPLOYEE, DEPARTMENT
```

- In this query, every EMPLOYEE is joined with every DEPARTMENT
- It is extremely important not to overlook specifying any selection and join conditions in the WHERE-clause; otherwise, incorrect and very large query results are produced

# Retrieving all the Attributes Using Asterisk (\*)

- To retrieve all the attribute values of the selected tuples, a \* (asterisk) is used, which stands for *all the attributes*

Examples:

```
Q1C:SELECT  *  
      FROM    EMPLOYEE  
      WHERE   DNO=5 ;
```

```
Q1D:SELECT  *  
      FROM    EMPLOYEE, DEPARTMENT  
      WHERE   DNAME='Research' AND  
              DNO=DNUMBER ;
```

# Eliminating Duplicates Using DISTINCT

- As mentioned earlier, SQL does not treat a relation as a set but a multiset (or bag); duplicate tuples can appear
- To eliminate duplicate tuples in a query result, the keyword **DISTINCT** is used
- Example: Result of Q11 may have duplicate SALARY values; result of Q11A does not have any duplicate values

```
Q11: SELECT  SALARY  
      FROM    EMPLOYEE
```

```
Q11A: SELECT DISTINCT SALARY  
      FROM    EMPLOYEE
```

# Set and Multiset Operations in SQL

- SQL has directly incorporated some set operations
- The set operations are: union (UNION), set difference (EXCEPT or MINUS) and intersection (INTERSECT)
- Results of these set operations are sets of tuples;  
*duplicate tuples are eliminated from the result*
- Set operations apply only to *type compatible* relations (also called *union compatible*); the two relations must have the same attribute types and in the same order
- Set operations typically applied to the results of two separate queries; e.g Q1 UNION Q2



# Set Operations (cont.)

- Example: Query 4: Make a list of all project names for projects that involve an employee whose last name is 'Smith' as a worker on the project or as a manager of the department that controls the project.

```
Q4:  (SELECT  PNAME
      FROM    PROJECT, DEPARTMENT, EMPLOYEE
      WHERE   DNUM=DNUMBER AND
MGRSSN=SSN AND LNAME='Smith')
      UNION
      (SELECT  PNAME
      FROM    PROJECT, WORKS_ON, EMPLOYEE
      WHERE   PNUMBER=PNO AND
              ESSN=SSN AND LNAME='Smith') ;
```

# Multiset Operations in SQL

- SQL has multiset operations when the user does not want to eliminate duplicates from the query results
- These are: UNION ALL, EXCEPT ALL, and INTERSECT ALL; see examples in Figure 4.5 (next slide)
- Results of these operations are multisets of tuples; *all tuples and duplicates in the input tables are considered* when computing the result
- Multiset operations also apply only to *type compatible* relations
- Typically applied to the results of two separate queries; e.g Q1 UNION ALL Q2 ;

(a)		(b)		(c)	
R		S		T	
A		A		A	
a1		a1		a2	
a2		a2		a3	
a2		a4			
a3		a5			

(b)		(c)	
T		T	
A		A	
a1		a2	
a1		a3	
a2			
a2			
a3			
a4			
a5			

(d)	
T	
A	
a1	
a2	

**Figure 4.5**

The results of SQL multiset operations. (a) Two tables, R(A) and S(A). (b) R(A) UNION ALL S(A). (c) R(A) EXCEPT ALL S(A). (d) R(A) INTERSECT ALL S(A).

# Substring Comparison Conditions

- The **LIKE** comparison operator is used to compare partial strings
- Two reserved characters are used: '\*' (or '%' in some implementations) replaces an arbitrary number of characters, and '\_' replaces a single arbitrary character
- Conditions can be used in WHERE-clause

# Substring Comparison Conditions (cont.)

- Example: Query 25: Retrieve all employees whose address is in Houston, Texas. Here, the value of the ADDRESS attribute must contain the substring 'Houston,TX' in it.

```
Q25: SELECT    FNAME, LNAME  
      FROM      EMPLOYEE  
      WHERE ADDRESS LIKE '*Houston,TX*' ;
```

# Substring Comparison Conditions (cont.)

- Example: Query 26: Retrieve all employees who were born during the 1950s.
  - Here, '5' must be the 3<sup>rd</sup> character of the string (according to the standard format for DATE yyyy-mm-dd), so the BDATE value is '\_\_5\_\_\_\_\_', with each underscore as a place holder for a single arbitrary character.

```
Q26:  SELECT  FNAME, LNAME
      FROM    EMPLOYEE
      WHERE   BDATE LIKE  '__5_____';
```

- The LIKE operator allows users to get around the fact that each value is considered atomic and indivisible
  - Hence, in SQL, character string attribute values are not atomic

# Applying Arithmetic in SQL Queries

- The standard arithmetic operators '+', '-', '\*', and '/' (for addition, subtraction, multiplication, and division, respectively) can be applied to numeric attributes and values in an SQL query
- Example: Query 27: Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
Q27:  SELECT  FNAME, LNAME, 1.1*SALARY
        FROM    EMPLOYEE, WORKS_ON, PROJECT
        WHERE   SSN=ESSN AND PNO=PNUMBER
                AND PNAME='ProductX' ;
```

# Ordering a query result using the ORDER BY clause

- The **ORDER BY** clause is used to sort the tuples in a query result based on the values of some attribute(s)
- Example: Query 28: Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.

```
Q28:  SELECT  DNAME, LNAME, FNAME, PNAME
        FROM    DEPARTMENT, EMPLOYEE,
                WORKS_ON, PROJECT
        WHERE   DNUMBER=DNO AND SSN=ESSN
                AND PNO=PNUMBER
        ORDER BY  DNAME, LNAME ;
```



# ORDER BY (cont.)

- The default order is in ascending order of values
- We can specify the keyword **DESC** if we want a descending order; the keyword **ASC** can be used to explicitly specify ascending order, even though it is the default
- Without ORDER BY, the rows in a query result appear in some system-determined (random) order

# Summary of Basic SQL Queries

- A basic query in SQL can consist of up to four clauses, but only the first two, SELECT and FROM, are mandatory
- Two additional clauses (GROUP BY and HAVING) will be discussed in Chapter 5
- The four basic clauses are specified in the following order:

**SELECT**    <attribute list>

**FROM**       <table list>

**[WHERE**    <condition>]

**[ORDER BY** <attribute list>]

# SQL Statements for Updating the Database

- There are three SQL commands to modify the database: **INSERT**, **DELETE**, and **UPDATE**
- **INSERT** is used for adding one or more records to a table
- **DELETE** is for removing records
- **UPDATE** is for modifying existing records
- Some operations may be rejected if they violate integrity constraints; others may propagate additional updating automatically if specified in the database schema

# INSERT statement

- Used to add one or more tuples (records) to a relation (table)
- Values for the attributes should be listed in the same order as the attributes were specified in the **CREATE TABLE** command
- Attributes that have default values can be omitted in the new record(s)

# INSERT statement (cont.)

- Example:

```
U1:  INSERT INTO    EMPLOYEE
      VALUES ('Richard','K','Marini', '653298653', '30-DEC-52',
              '98 Oak Forest,Katy,TX', 'M', 37000,'987654321', 4 ) ;
```

- An alternate form of INSERT specifies explicitly the attribute names that correspond to the values in the new tuple
  - Attributes with NULL or default values can be left out
- Example: Insert a tuple for a new EMPLOYEE for whom we only know the FNAME, LNAME, and SSN attributes.

```
U1A:  INSERT INTO EMPLOYEE (FNAME, LNAME, SSN)
      VALUES ('Richard', 'Marini', '653298653') ;
```

# INSERT statement (cont.)

- Constraints specified in the DDL are automatically enforced by the DBMS when updates are applied to the database
- Can insert multiple tuples in one INSERT statement
  - The values in each tuple are enclosed in parentheses

# INSERT statement (cont.)

- Can also insert tuples from a query result into a table
- Example: Suppose we want to create a temporary table that has the employee last name, project name, and hours per week for each employee.
  - A table WORKS\_ON\_INFO is created by U3A, and is loaded with the summary information retrieved from the database by the query in U3B.

```
U3A:  CREATE TABLE WORKS_ON_INFO
      (EMP_NAME          VARCHAR(15),
       PROJ_NAME         VARCHAR(15),
       HOURS_PER_WEEK    DECIMAL(3,1));
```

```
U3B:  INSERT INTO  WORKS_ON_INFO (EMP_NAME,
                                PROJ_NAME, HOURS_PER_WEEK)
      SELECT      E.LNAME, P.PNAME, W.HOURS
      FROM        EMPLOYEE E, PROJECT P, WORKS_ON W
      WHERE       E.SSN=W.ESSN AND W.PNO=P.PNUMBER ;
```

# INSERT statement (cont.)

- Note: The WORKS\_ON\_INFO table may not be up-to-date if we change the tuples in the WORKS\_ON, PROJECT, or EMPLOYEE relations *after* executing U3B.
- We have to use CREATE VIEW (see Chapter 5) to keep such a table up to date.



# DELETE Statement

- Removes tuples from a relation
  - Includes a WHERE-clause to select the tuples to be deleted
  - Referential integrity should be enforced (via REJECT, CASCADE, SET NULL, or SET DEFAULT)
  - Tuples are deleted from only *one table* at a time (unless CASCADE is specified on a referential integrity constraint)
  - Missing WHERE-clause deletes *all tuples* in the relation; the table then becomes an empty table
  - Number of tuples deleted is the number of tuples selected by the WHERE-clause

# DELETE Statement (cont.)

- Examples:

U4A:   DELETE FROM       EMPLOYEE  
          WHERE            LNAME='Brown' ;

U4B:   DELETE FROM       EMPLOYEE  
          WHERE            SSN='123456789' ;

U4C:   DELETE FROM       EMPLOYEE  
          WHERE            DNO = 5 ;

U4D:   DELETE FROM       EMPLOYEE ;

# UPDATE Statement

- Used to modify attribute values of one or more selected tuples
- A WHERE-clause selects the tuples to be modified
- An additional SET-clause specifies the attributes to be modified and their new values
- Each command modifies tuples *in the same relation*
- Referential integrity should be enforced (via REJECT, CASCADE, SET NULL, or SET DEFAULT)

# UPDATE Statement (cont.)

- Example: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

```
U5:  UPDATE  PROJECT  
      SET  PLOCATION = 'Bellaire', DNUM = 5  
      WHEREPNUMBER=10 ;
```

# UPDATE Statement (cont.)

- Example: Give all employees in department number 5 a 10% raise in salary  
U6:     UPDATE   EMPLOYEE  
       SET         SALARY = SALARY \*1.1  
       WHERE     DNO = 5 ;
- In this request, the modified SALARY value depends on the original SALARY value in each tuple
  - The reference to the SALARY attribute on the right of = refers to the old SALARY value before modification
  - The reference to the SALARY attribute on the left of = refers to the new SALARY value after modification

# Some Other Features of SQL

- More complex query features: nested queries, aggregate functions, GROUP BY, HAVING, EXISTS function (see Chapter 5)
- Views, assertions, schema modification (Chapter 5)
- Triggers (Chapter 5 and Section 26.1)
- Object-relational features (Chapter 11, Section 11.2)
- Programming techniques (Chapter 13)
- Transaction features (Chapter 21, Section 21.6)
- Database Security (Chapter 24, Section 24.2)

# Chapter 4 Summary

- Overview of SQL
- SQL Data Definition (DDL) for Specifying a Relational Database Schema
  - CREATE statements
  - Specifying Constraints
- Basic Retrieval Queries in SQL
  - SELECT ... FROM ... WHERE ... statements
- Basic Database Modification in SQL
  - INSERT, DELETE, UPDATE statements