

An introduction to
NoSQL databases

What is covered in this

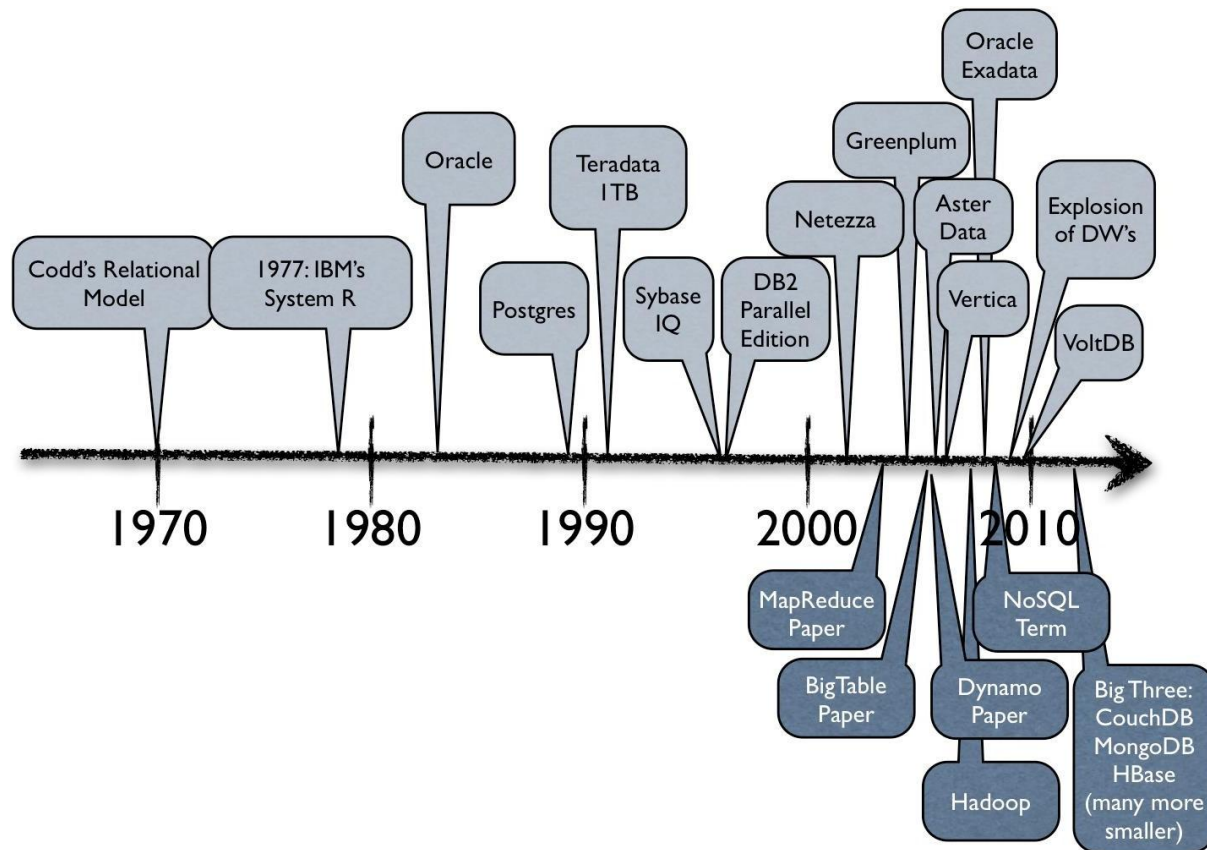
2

- A brief history of databases
- NoSQL WHY, WHAT & WHEN?
- Characteristics of NoSQL databases
- Aggregate data models
- CAP theorem

- ❑ Database - Organized collection of data
- ❑ DBMS Database Management System: a package with software and computer programs that controls the creation, maintenance and use of a database
- ❑ Databases are created to operate large quantities of information by inputting, storing, retrieving, and managing that information.

A brief history

4



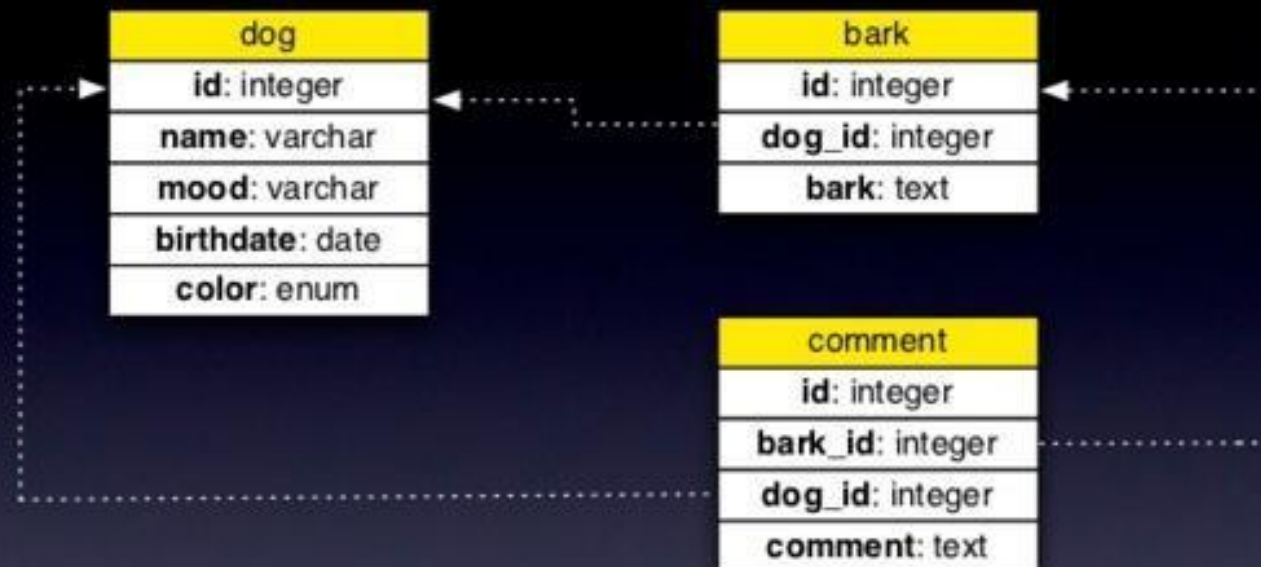
- Benefits of Relational databases:

- Designed for all purposes
- ACID
- Strong consistancy, concurrency, recovery
- Mathematical background
- Standard Query language (SQL)
- Lots of tools to use with i.e: Reporting services, entity frameworks, ...

SQL databases

6





id	name	mood	birth_date	color
12	Stella	Happy	2007-04-01	NULL
13	Wimma	Hungry	NULL	black
9	Ninja	NULL	NULL	NULL

NoSQL why, what and when?

8


But...

- ❑ Relational databases were not built for **distributed applications**.

Because...

- ❑ Joins are expensive
- ❑ Hard to scale horizontally
- ❑ Impedance mismatch occurs
- ❑ Expensive (product cost, hardware, Maintenance)

Era of Distributed Computing



NoSQL why, what and when?

9

But...

- ☐ Relational databases were not built for **distributed applications**.


Because...

- ☐ Joins are expensive
- ☐ Hard to scale horizontally
- ☐ Impedance mismatch occurs
- ☐ Expensive (product cost, hardware, Maintenance)

And....

It's weak in:

- ☐ Speed (performance)
- ☐ High availability
- ☐ Partition tolerance




Why NOSQL now??

Ans. Driving

11

New Trends

Trends



Big data



Connectivity



P2P Knowledge



Concurrency



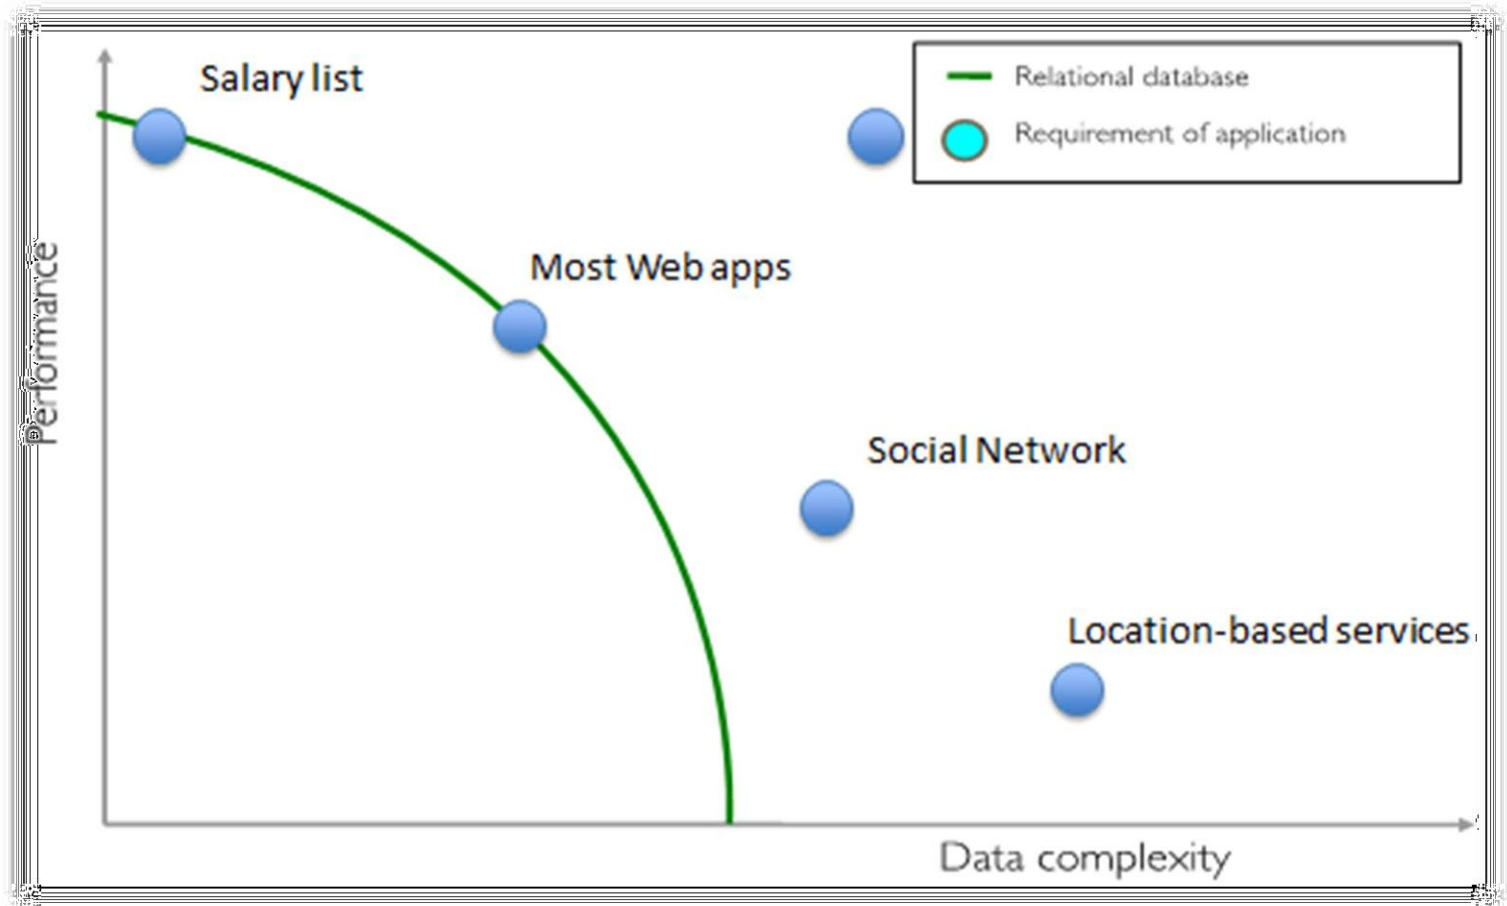
Diversity



Cloud-Grid

Side note: RDBMS

12



But... What's

- ❑ **No SQL?** A No SQL database provides a mechanism for storage and retrieval of data that employs less constrained consistency models than traditional relational database
- ❑ No SQL systems are also referred to as "NotonlySQL" to emphasize that they do in fact allow SQL-like query languages to be used.



Characteristics of NoSQL databases


14

NoSQL avoids:

- ❑ Overhead of ACID transactions
- ❑ Complexity of SQL query
- ❑ Burden of up-front schema design
- ❑ DBA presence
- ❑ Transactions (It should be handled at application layer)

Provides:

- ❑ Easy and frequent changes to DB
- ❑ Fast development
- ❑ Large data volumes(eg.Google)
- ❑ Schema less



NoSQL why, what and when?

10

When and when not to use it?

WHEN / WHY ?

- When traditional RDBMS model is too restrictive (flexible schema)
- When ACID support is not "really" needed
- Object-to-Relational (O/R) impedance
- Because RDBMS is neither distributed nor scalable by nature
- Logging data from distributed sources
- Storing Events / temporal data
- Temporary Data (Shopping Carts / Wish lists / Session Data)
- Data which requires flexible schema
- **Polyglot Persistence** i.e. best data store depending on nature of data.

WHEN NOT ?

- Financial Data
- Data requiring strict ACID compliance
- Business Critical Data

NoSQL is getting more & more popular

15

Google

ebay™


LinkedIn™

YAHOO!®

NETFLIX

amazon

theguardian



facebook®

What is a schema-less datamodel?

16

In relational Databases:

- ❑ You can't add a record which does not fit the schema
- ❑ You need to add NULLs to unused items in a row
- ❑ We should consider the datatypes.
i.e : you can't add a string to an integer field
- ❑ You can't add multiple items in a field (You should create another table: primary-key, foreign key, joins, normalization, ... !!!)

```
create table customers (id int, firstname text, lastname text)  
insert into customers (firstname, middlename, lastname) values (...)
```




What is a schema-less datamodel?

17


In NoSQL Databases:

- There is no schema to consider
- There is no unused cell
- There is no datatype (implicit)
- Most of considerations are done in application layer
- We gather all items in an aggregate (document)

Relational Model



Document Model



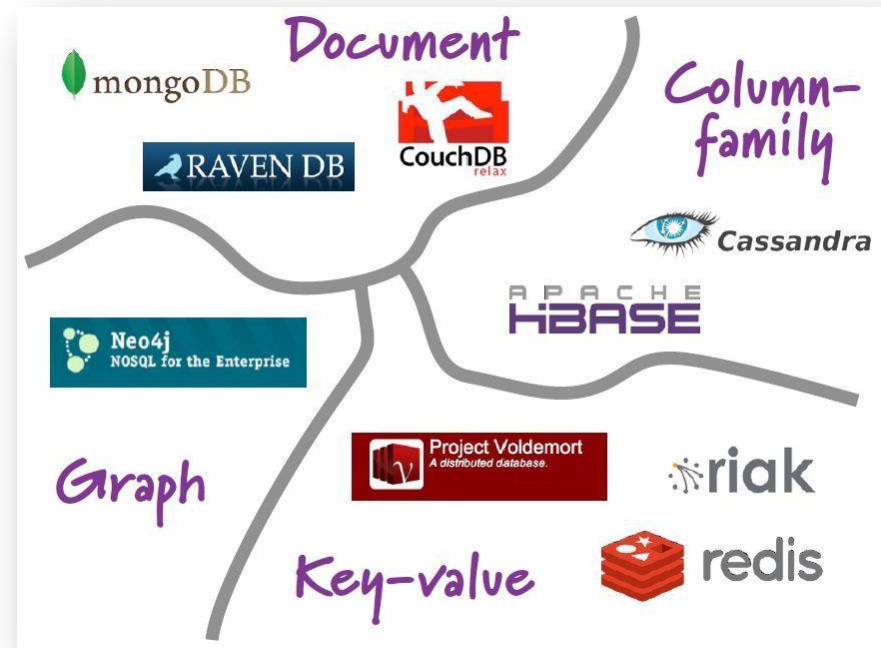
Aggregate Data Models

18

NoSQL databases are classified in four major datamodels:

- Key-value
- Document
- Column family
- Graph

Each DB has its own query language

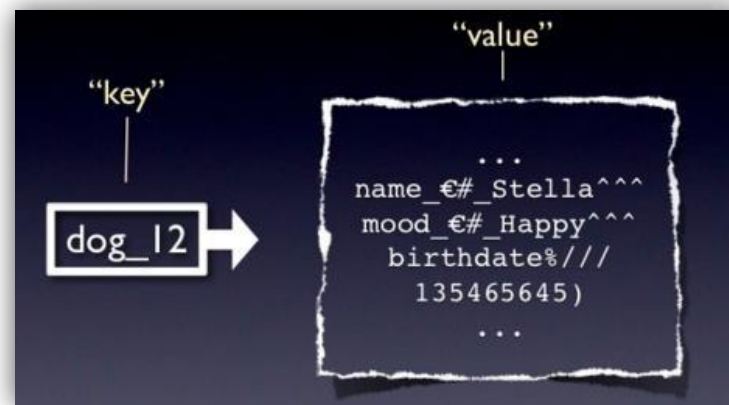


Key-value data model

19

- Simplest NOSQL databases
- The main idea is the use of a hash table
- Access data (values) by strings called keys
- Data has no required format data may have any format
- Data model: (key, value) pairs
- Basic Operations:
 - Insert(key,value),
 - Fetch(key),
 - Update(key),
 - Delete(key)

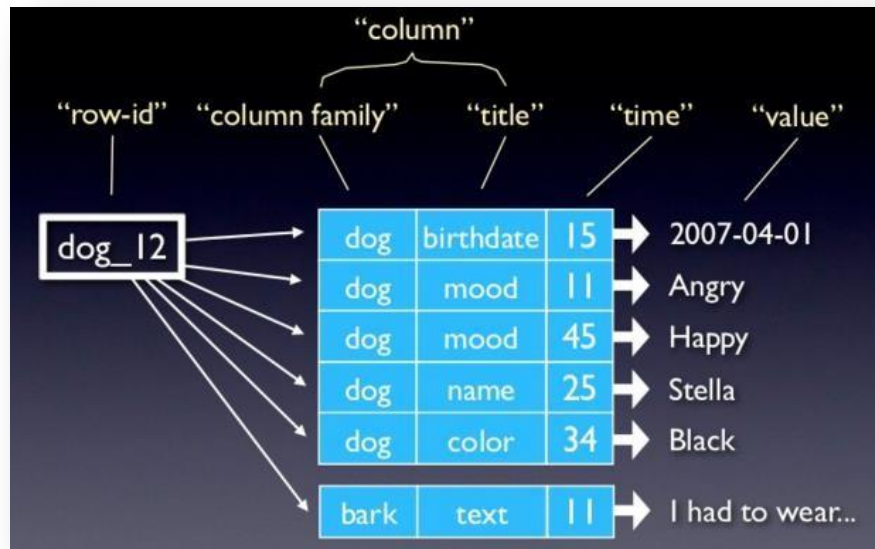
Car	
Key	Attributes
1	Make: Nissan Model: Pathfinder Color: Green Year: 2003
2	Make: Nissan Model: Pathfinder Color: Blue Color: Green Year: 2005 Transmission: Auto



Column family data model

20

- The column is lowest/smallest instance of data.
- It is a tuple that contains a name, a value and a timestamp



ColumnFamily: Authors		
Key	Value	
"Eric Long"	Columns	
	Name	Value
	"email"	"eric (at) long.com"
	"country"	"United Kingdom"
	"registeredSince"	"01/01/2002"
"John Steward"	Columns	
	Name	Value
	"email"	"john.steward (at) somedomain.com"
	"country"	"Australia"
	"registeredSince"	"01/01/2009"
"Ronald Mathies"	Columns	
	Name	Value
	"email"	"ronald (at) sodeso.nl"
	"country"	"Netherlands, The"
	"registeredSince"	"01/01/2010"

Some statistics about Facebook Search (using **Cassandra**)

- ❖ MySQL > 50 GB Data
 - Writes Average : ~300 ms
 - Reads Average : ~350 ms

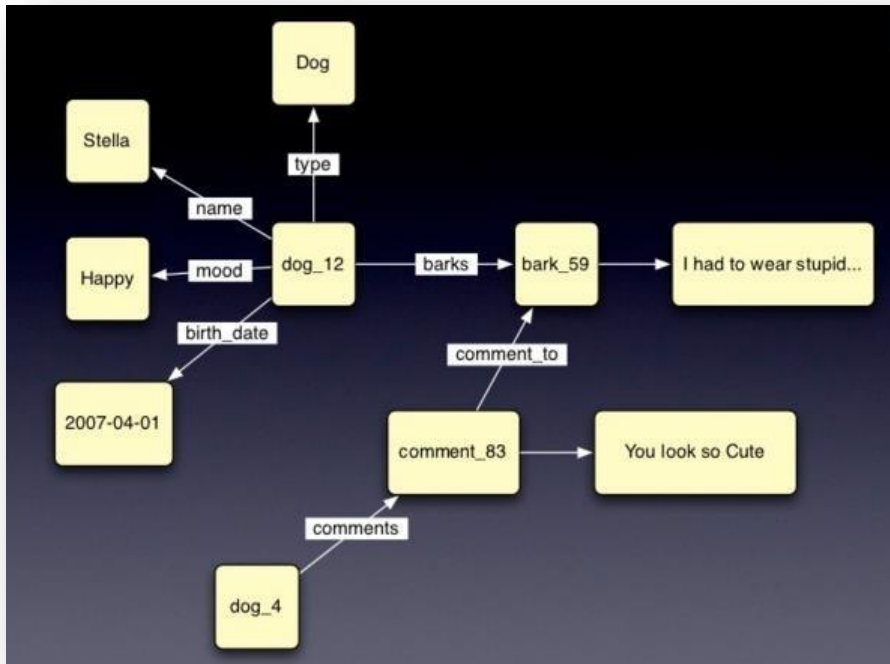
- ❖ Rewritten with Cassandra > 50 GB Data
 - Writes Average : 0.12 ms
 - Reads Average : 15 ms



Graph data model

22

- Based on Graph Theory.
- Scale vertically, no clustering.
- You can use graph algorithms easily
- Transactions
- ACID

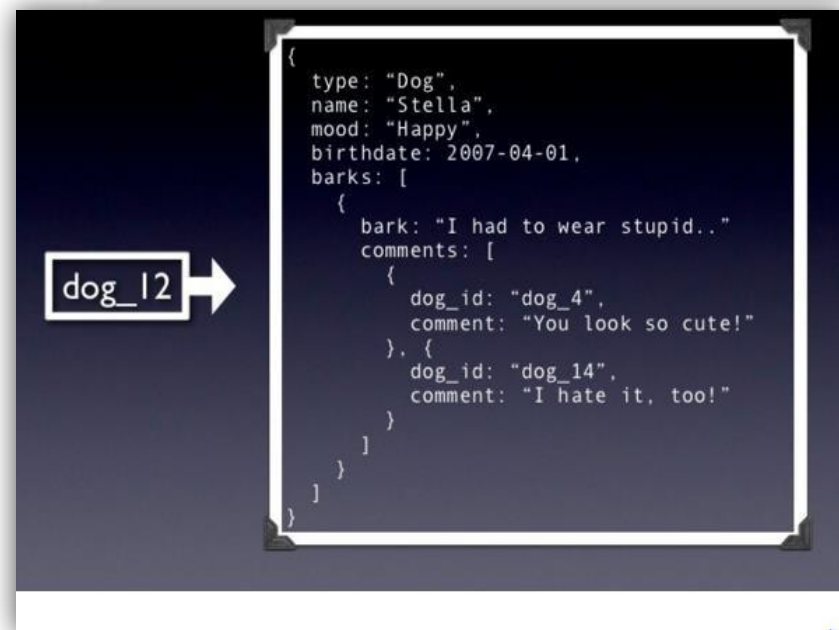
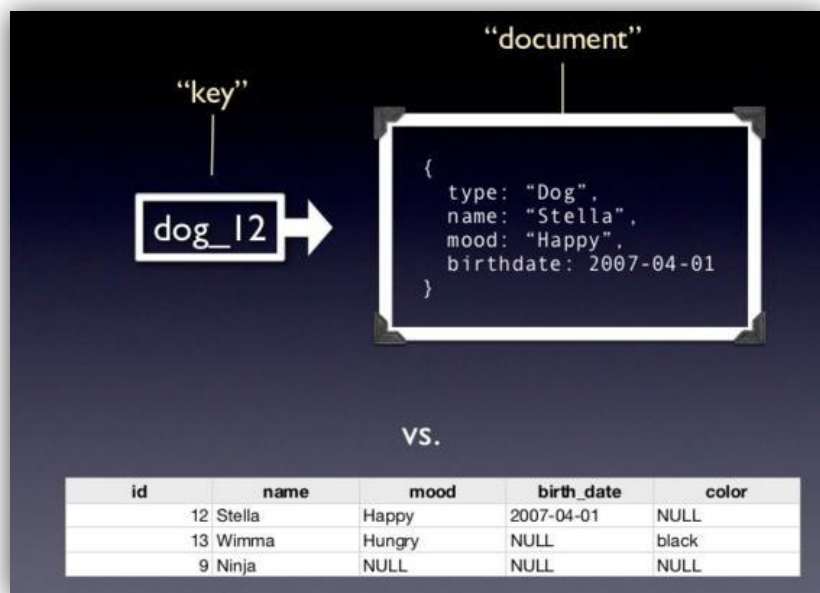


Document based data model

23

- Pair each key with complex data structure known as data structure.
- Indexes are done via B-Trees.
- Documents can contain many different key-value pairs, or key-array pairs, or even nested documents.

```
{
  person: {
    first_name: "Peter",
    last_name: "Peterson",
    addresses: [
      {street: "123 Peter St"},
      {street: "504 Not Peter St"}
    ],
  },
}
```



Document based data model

24

```
SELECT name, pic, profile_url
FROM user
WHERE uid = me()
```



```
SELECT message, attachment
FROM stream
WHERE source_id = me() AND type = 80
```

```
SELECT name
FROM friendlist
WHERE owner = me()
```



```
SELECT name
FROM group
WHERE gid IN ( SELECT gid
                FROM group_member
                WHERE uid = me() )
```



```
SELECT name, pic
FROM user
WHERE online_presence = "active"
AND uid IN ( SELECT uid2
              FROM friend
              WHERE uid1 = me() )
```



<https://developers.facebook.com/docs/reference/fql/>

Differences

	SQL Databases	No SQL Database
Example	Oracle , mysql	Mondo DB, CouchDB, Neo4J
Storage Model	Rows and tables	Key-value. Data stored as single document in JSON, XML
Schemas	Static	Dynamic
Scaling	Vertical & Horizontal	Horizontal
Transactions	Yes	Certain levels
Data Manipulation	Select, Insert , Update	Through Object Oriented API's


- We need a distributed database system having such features:
 - **Fault tolerance**
- - **High availability**
- - **Consistency**
- - **Scalability**
-

Which is impossible!!!

According to CAP theorem

The CAP Theorem

- Impossible for any shared data-system to guarantee simultaneously all of the following three properties:
 - Consistency – once data is written, all future read requests will contain that data
 - Availability – the database is always available and responsive
 - Partition Tolerance – if part of the database is unavailable, other parts are unaffected

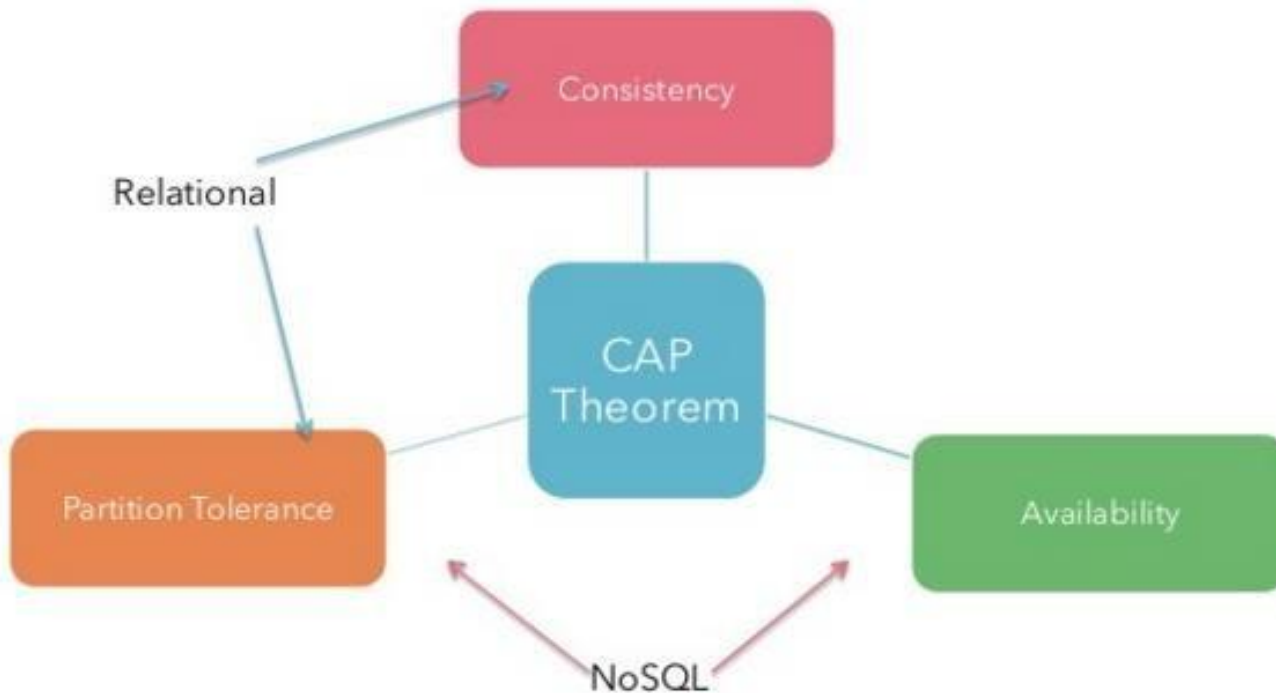


We can not achieve all the three items
In distributed database systems
(center)

CAP theorem

28

Traditional vs. NoSQL



In Conclusion!

- RDBMS is a great tool for solving ACID problems
 - When data validity is super important
 - When you need to support dynamic queries
- NoSQL is a great tool for solving data availability problems
 - When it's more important to have fast data than right data
 - When you need to scale based on changing requirements
- Pick the right tool for the job

- ❑ nosql-database.org/
- ❑ <https://www.mongodb.com/nosql-explained>
- ❑ www.couchbase.com/nosql-resources/what-is-no-sql
- ❑ <http://nosql-database.org/> "NoSQL DEFINITION: Next Generation Databases mostly addressing some of the points: being non-relational, distributed, open-source and horizontally scalable"
- ❑ NoSQL distilled, Martin Fowler
- ❑ Please like and follow at www.slideshare.net/AshwaniKumar274

Thanks...

Any Questions??